

A dark blue vertical bar on the left side of the page. A blue arrow points to the right from the bar, containing the text 'BTP-II'.

BTP-II

AIRCRAFT DETECTION IN SATELLITE IMAGES

TESTING & DEPLOYMENT

Several thin, curved lines in shades of blue and grey originate from the bottom left and curve upwards and to the right.

Geordi Jose

DEPT. OF CIVIL ENGINEERING
IIT BOMBAY

ABSTRACT

With the previous BTP-I carried out on detecting whether an image contains an aircraft or not and simultaneously another methodology applied to find out where in a given image aircraft is present, with this BTP I aimed at creating a merged and faster solution which can detect and locate (using bounding boxes) aircrafts in digital satellite images. The problem was divided into two problems, with my partner working on preparing a deep learning object detection model, and with me working on testing, validating and deploying the said model. Hyperparameters of the model like epochs, iterations per epoch were improved by testing on dataset and feeding back the results. The mAP (mean Average Precision) score is used to evaluate the model. After testing, a deployable version of our solution was built using ANVIL, an open source tool for developing full stack python web applications. A simple front-end UI was built, which interacts with a python server script on the backend where any user can simply upload an image and run the model on it. The results obtained are promising on the given set of testing images, but generalizes very poorly outside of it. This is due to a very limited artificially generated training dataset and therefore there is much scope in for this project if a larger more versatile dataset can be used.

INTRODUCTION

Object detection in images has been area of research for over many decades and among these, aircraft detection still remains a challenging task because of the complex background, illumination change and variation of aircraft kind. The detection still needs to be carried out for military and security purpose. Various different approach has been carried out for aircraft detection starting from using filter with support vector or coarse-to-fine edge detection, using directional gradient and so on. But these methods use low-level features and thus has high false alarm. With the introduction of neural networks works shifted towards using supervised training approach and deep learning methods. Running deep networks on sliding windows for object proposal proved to be accurate than previous approaches but was very slow in execution. To increase the speed, detection based on spatial pyramid pooling method was adopted but the issue of slow candidate region proposal network still remained. Later the concept of using convolutional neural network along with region proposal network was introduced which proved to be very accurate and also reducing the runtime of the execution. The key requirements of using a convolutional structure is that as it involves so many weight parameters so the training data set needs to be very accurate and present in large number. To compensate for large training data set, data augmentation becomes a crucial aspect in this method. Also, the hidden structure of CNN doesn't follow a fixed pattern so finding the appropriate CNN hidden layers varies from the targeted purpose and the kind of dataset that we have. In this project Faster-RCNN technique with region proposal network was incorporated to cater to detection and creating bounding box along the object.

Evaluating the accuracy of the model is also a challenging task as it is not quite as simple as a classification problem. AP (Average precision) is a popular metric in measuring the accuracy of object detectors like Faster R-CNN, SSD, etc. The general definition for the

Average Precision (AP) is finding the area under the precision-recall curve. In our case we have an Intersection over Union (IoU) >0.5 criterion to define a true positive (TP) case with respect to accuracy. With the TP, FP and FN formally defined in this manner, we can now calculate the precision and recall of our detection for a given class across the test set and use this for calculating average precision.

An aircraft detection solution is expected to have several applications in defence/military as well as civil aviation sectors. Therefore it is important to deploy this work accessible to any ordinary user. Anvil was used to create a web app front end for this work, which interacts with users to accept input images and display output images.

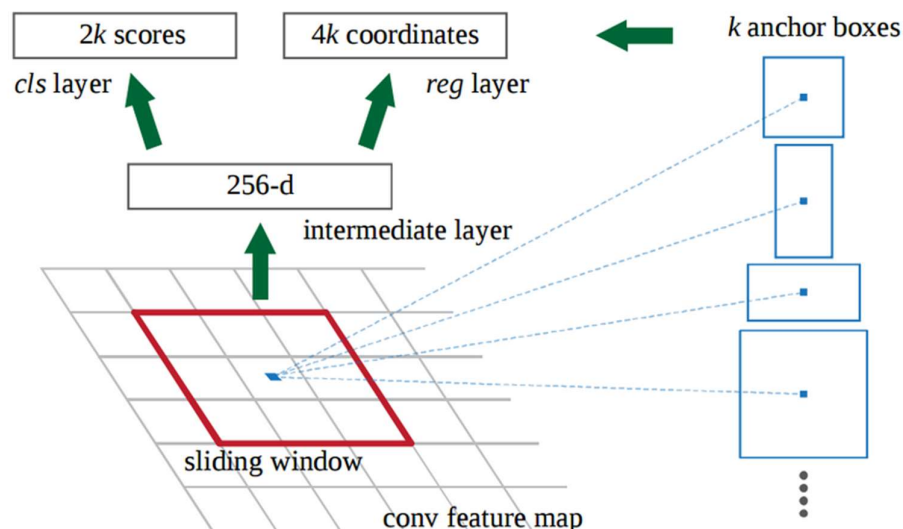
BACKGROUND

BRIEF EXPLANATION OF FASTER- RCNN MODEL EMPLOYED

The main insight of Faster R-CNN was to replace the slow selective search algorithm of Fast RCNN with a fast-neural net. Specifically, it introduced the region proposal network (RPN). Here's how the RPN works:

- At the last layer of an initial CNN, a 3×3 sliding window moves across the feature map and maps it to a lower dimension (e.g. 256-d)
- For each sliding-window location, it generates multiple possible regions based on k fixed-ratio anchor boxes (default bounding boxes)
- Each region proposal consists of a) an "objectness" score for that region and b) 4 coordinates representing the bounding box of the region

In other words, we look at each location in our last feature map and consider k different boxes centered around it: a tall box, a wide box, a large box, etc. For each of those boxes, we output whether or not we think it contains an object, and what the coordinates for that box are. This is what it looks like at one sliding window location:



The 2k scores represent the softmax probability of each of the k bounding boxes being on “object.” Notice that although the RPN outputs bounding box coordinates, it does not try to classify any potential objects: its sole job is still proposing object regions. If an anchor box has an “objectness” score above a certain threshold, that box’s coordinates get passed forward as a region proposal.

Once we have our region proposals, we feed them straight into what is essentially a Fast R-CNN. We add a pooling layer, some fully-connected layers, and finally a softmax classification layer and bounding box regressor. In a sense, Faster R-CNN = RPN + Fast R-CNN.

BRIEF EXPLANATION OF MEAN AVERAGE PRECISION (mAP) SCORE

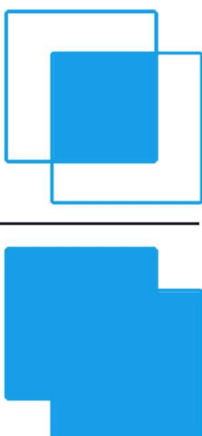
Precision and recall are two commonly used metrics to judge the performance of a given classification model. To understand mAP, we would need to first understand precision and recall. Precision of a given class in classification, a.k.a. positive predicted value, is given as the ratio of true positive (TP) and the total number of predicted positives.

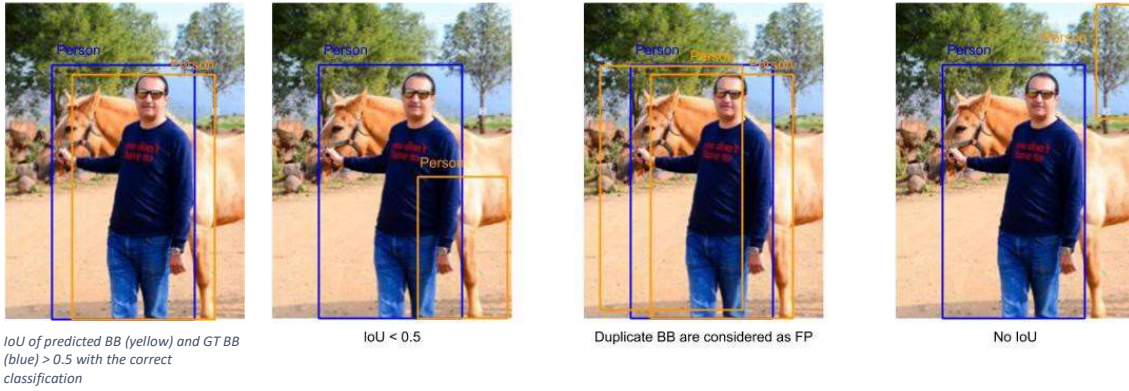
$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

The recall, a.k.a. true positive rate or sensitivity, of a given class in classification, is defined as the ratio of TP and total of ground truth positives.

$$\text{Recall (TPR)} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$


Obviously, there lies a trade-off between its precision and recall performance. For our precision to be high, we would need to decrease our number of FP, by doing so, it will decrease our recall. Similarly, by decreasing our number of FN would increase our recall and decrease our precision. Very often for information retrieval and object detection cases, we would want our precision to be high (our predicted positives to be TP). If we are using a neural network, this trade-off can be adjusted by the model’s final layer softmax threshold. But to apply this to an object detection model, we need to define what a True positive or False Negative case is. For this we use the concept of Intersection over Union (IoU). The IoU is given by the ratio of the area of intersection and area of union of the predicted bounding box and ground truth bounding box.

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$




Traditionally, we define a prediction to be a TP if the IoU is > 0.5. There are two possible scenarios where a BB would be considered as FP: an IoU < 0.5 or a Duplicated Bounding Box. For our single class detection problem, False Negative is when there is no detection at all. With the TP, FP and FN formally defined, we can now calculate the precision and recall of our detection for a given class across the test set. The general definition for the Average Precision (AP) is finding the area under the precision-recall curve.

$$AP = \int_0^1 p(r)dr$$



Precision and recall are always between 0 and 1. Therefore, AP falls within 0 and 1 also. The mAP for object detection is the average of the AP calculated for all the classes, which in our case only across the airplane class.

METHODOLOGY

DATASET

Dataset for testing and training included a set of 500 computer generated imagery (CGI) of satellite images with aircrafts placed randomly across the image. 400 images were used for augmentation and training followed by testing on 100 test images.



Sample Test images

train_labels.csv - Excel

File Home Insert Page Layout Formulas Data Review View Help Tell me what you want to do

Paste Cut Copy Format Painter Clipboard Font Alignment Number Conditional Formatting Table

A1 filename

	A	B	C	D	E	F	G	H	I	J	K	L	M
1	filename	xmin	ymin	xmax	ymax	label							
2	img1.png	410	356	469	419	plane							
3	img1.png	635	276	738	380	plane							
4	img1.png	830	46	902	119	plane							
5	img1.png	669	80	715	120	plane							
6	img1.png	58	225	108	272	plane							
7	img1.png	526	279	579	332	plane							
8	img1.png	822	446	870	493	plane							
9	img1.png	591	322	664	400	plane							
10	img1.png	838	227	920	309	plane							
11	img1.png	605	353	661	415	plane							
12	img1.png	775	434	827	484	plane							
13	img1.png	159	467	235	541	plane							
14	img10.png	439	537	475	574	plane							
15	img10.png	156	277	251	376	plane							
16	img10.png	480	480	524	526	plane							
17	img10.png	763	39	860	136	plane							
18	img10.png	50	318	123	392	plane							
19	img10.png	568	324	621	376	plane							
20	img10.png	375	134	420	182	plane							
21	img10.png	490	486	531	528	plane							
22	img10.png	245	493	317	567	plane							
23	img10.png	859	59	909	111	plane							
24	img10.png	839	152	905	221	plane							
25	img100.pn	537	69	595	127	plane							
26	img100.pn	307	203	343	241	plane							
27	img100.pn	820	162	888	229	plane							
28	img100.pn	490	97	586	190	plane							
29	img100.pn	118	30	193	104	plane							

train_labels

Ready

Bounding box coordinates for each image

TESTING AND mAP SCORE CALCULATION

Testing was performed in Google Colab GPU runtimes, for efficient results, the same environment where the model was trained. Each image is passed through the RPN network followed by the trained CNN. Hyperparameters for the model like the number of epochs, iterations per epoch, threshold values were fed back into the model after testing based on mAP scores obtained. An accuracy threshold of 0.9 is set for classification i.e the model will locate and display an aircraft only if the final associated probability of this being an aircraft is greater than 0.9 (or 90 percent)

```
def get_map(pred, gt, F):
    T = {}
    P = {}
    Fx, Fy = F

    for bbox in gt:
        bbox['bbox_matched'] = False

    pred_probs = np.array([s['prob'] for s in pred])
    box_ids_sorted_by_prob = np.argsort(pred_probs)[::-1]

    for box_idx in box_ids_sorted_by_prob:
        pred_box = pred[box_idx]
        pred_class = pred_box['class']
        pred_x1 = pred_box['x1']
        pred_x2 = pred_box['x2']
        pred_y1 = pred_box['y1']
        pred_y2 = pred_box['y2']
        pred_prob = pred_box['prob']

        if pred_class not in P:
            P[pred_class] = []
            T[pred_class] = []
            P[pred_class].append(pred_prob)
            found_match = False

        for gt_box in gt:
            gt_class = gt_box['class']
            gt_x1 = gt_box['x1']/Fx
            gt_x2 = gt_box['x2']/Fx
            gt_y1 = gt_box['y1']/Fy
            gt_y2 = gt_box['y2']/Fy
            gt_seen = gt_box['bbox_matched']
            if gt_class != pred_class:
                continue
            if gt_seen:
                continue
            iou_map = iou((pred_x1, pred_y1, pred_x2, pred_y2), (gt_x1, gt_y1, gt_x2, gt_y2))
            if iou_map >= 0.5:
                found_match = True
                gt_box['bbox_matched'] = True
                break
            else:
                continue

        T[pred_class].append(iou(found_match))

    for gt_box in gt:
        if not gt_box['bbox_matched'] and not gt_box['difficult']:
            if gt_box['class'] not in P:
                P[gt_box['class']] = []
                T[gt_box['class']] = []
            T[gt_box['class']].append(1)
            P[gt_box['class']].append(0)

    import pdb
    pdb.set_trace()
    return T, P
```

Code for mAP score computation

```
Elapsed time = 0.55113101005542
plane AP: 0.5401217304735948
mAP = 0.5401217304735948
94/100
Elapsed time = 0.5796689987182617
plane AP: 0.5397114055931123
mAP = 0.5397114055931123
95/100
Elapsed time = 0.5873677730560303
plane AP: 0.5324633435070996
mAP = 0.5324633435070996
96/100
Elapsed time = 0.567589521408081
plane AP: 0.5299526700697973
mAP = 0.5299526700697973
97/100
Elapsed time = 0.5438728332519531
plane AP: 0.5290003837731405
mAP = 0.5290003837731405
98/100
Elapsed time = 0.5564498901367188
plane AP: 0.527798500775934
mAP = 0.527798500775934
99/100
Elapsed time = 0.5385639667510986
plane AP: 0.5269807580818849
mAP = 0.5269807580818849

mean average precision: 0.5386784240274048
```

mAP scores for each image followed by average

WEB APP FOR DEPLOYING THE MODEL

The front end web app was created using a free open source tool ANVIL, which enables developing full webpages using python alone. This web app enables user to upload an input image to the model and view results without any technical hindrance. If necessary, users can also specify the accuracy threshold they want to be applied for object detection.

```
from ..anvil_designer import Form1Template
from anvil import *
import anvil.server

class Form1(Form1Template):

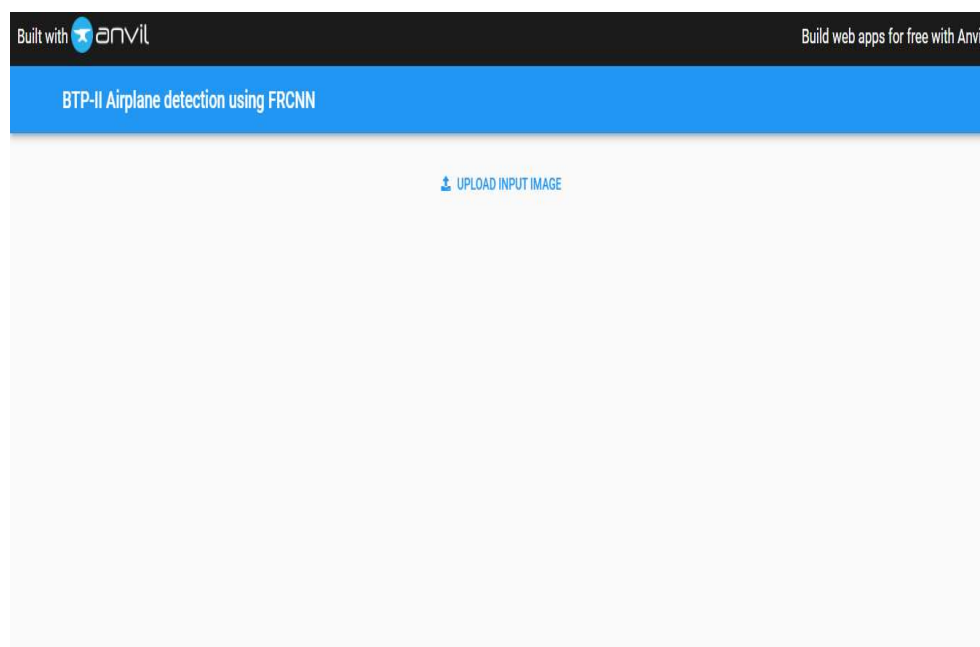
    def __init__(self, **properties):
        # Set Form properties and Data Bindings.
        self.init_components(**properties)
        self.label_1.text= "BTP-II Airplane detection using FRCNN"
        self.file_loader_1.text="Upload Input Image "
        self.image_1.display_mode='original_size'
        self.image_2.display_mode='original_size'

        # Any code you write here will run when the form opens.

    def file_loader_1_change(self, file, **event_args):
        """This method is called when a new file is loaded into this FileLoader"""
        self.image_1.source=self.file_loader_1.file
        my_media=self.file_loader_1.file
        processed_image=anvil.server.call('process_image', my_media)
        self.image_2.source=processed_image
```

The anvil code's file loader change function is called when the user inputs an image. Subsequently the image is passed on to a server side script by invoking the 'process-image' function defined there.

ANVILcode



Anvil Design


```

[ ] Connecting to wss://anvil.works/uplink
Anvil websocket open
[ ] Authenticated OK

@anvil.server.callable
def process_image(my_media):

    import numpy as np
    import cv2

    # read in image from media object
    arr = np.fromstring(my_media.get_bytes(), np.uint8)
    im = cv2.imdecode(arr, cv2.IMREAD_COLOR)
    cv2.imwrite('input.png', im)
    #####
    ifilepath='input.png'
    findplane(ifilepath)
    ofilepath=('result.png')
    im=cv2.imread(ofilepath)
    #####

    # convert image to string and return downloadable media
    im_str = cv2.imencode('.png', im)[1].tostring()

    return anvil.BlobMedia("image/png", im_str, name='myimage.png')


```



Server Side Script on Colab

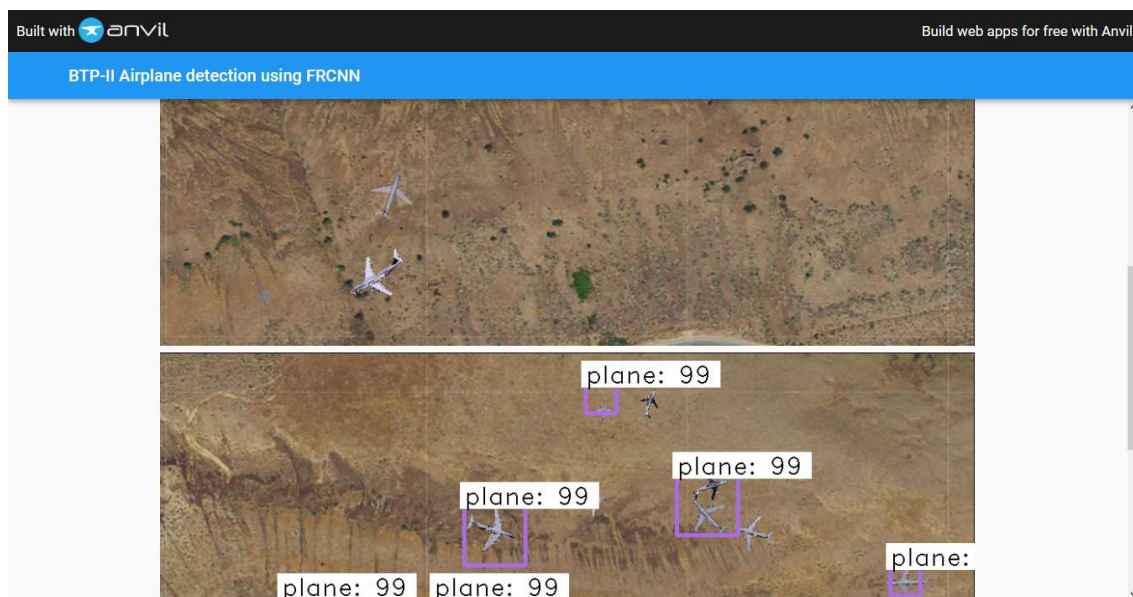

The server side script on Google Colab. The process image function receives the media object from anvil, and after converting this to an image file, the findplane function is called which predicts airplanes in the image using the trained and tested FRCNN image. The output image of this function is passed back to anvil after converting it to downloadable media.

RESULTS

After final testing an mAP score of 0.54 was achieved. Some of the test result images:







1ANVIL web app for airplane detection

CONCLUSIONS

A solution has been developed to identify and locate aircrafts from digital images using a Faster R-CNN deep learning model in the backend. The model is developed using artificially generated satellite imagery. The model performs really well on images similar to the dataset it was trained on, and achieved an accuracy score (mAP) of 0.54 after testing on 100 images. In the limited domain the model has been trained on, it performs really well, and therefore proves promising for working on larger datasets. As of now performance outside the trained dataset is poor i.e the model does not generalise well. A web app front end for accessing the model has also been developed.

LIMITATIONS AND SCOPE FOR FUTURE WORK

With the objective of locating any aircraft present in an image, this model has its own restrictions. Beyond the realms of training dataset, it generalizes very poorly and there are two major reason for this. First being the original dataset, itself is being computer generated and not comes from original surrounding. The aircraft has been put randomly on different backgrounds which leads to low representation of actual environment. Second major reason is the limited amount of dataset the model is trained on. There were only 400 original images on which augmentation was performed to increase it to 4400. Having a large set of original datasets covering all general environments would certainly boost the accuracy.

For the usage of this model, a UI has been hosted on anvil server where any use can upload any image and then would get the output with bounding box created around aircrafts present if any in that image. If we have a capable enough backend server we can even run this on a video input, splitting each video into multiple frames and then performing testing on each generated image.

As the entire framework has been developed such that it can detect any given object provided sufficient training has been done, so expanding the horizon of the dataset will lead to model predicting many different types of objects as and when trained upon. With bigger more versatile datasets, rather than looking at a single aircraft class, looking for multiple aircraft classes i.e identifying even the type of aircrafts is achievable.

REFERENCES

1. Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks [Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun] Microsoft Research{v-shren, kahe, rbg, jiansun}@microsoft.com
2. X. Li, S. Wang, B. Jiang and X. Chan, "Airplane detection using convolutional neural networks in a coarse-to-fine manner," 2017 IEEE 2nd Information Technology, Networking, Electronic and Automation Control Conference (ITNEC), Chengdu, 2017, pp. 235-239.
doi:10.1109/ITNEC.2017.8284943

3. Xu, Yuelei et al. "Rapid Airplane Detection in Remote Sensing Images Based on Multilayer Feature Fusion in Fully Convolutional Neural Networks." *Sensors (Basel, Switzerland)* vol. 18,7 2335. 18 Jul. 2018, doi:10.3390/s18072335
4. Chen, Xueyun, Shiming Xiang, Cheng-Lin Liu and Chunhong Pan. "Aircraft Detection by Deep Convolutional Neural Networks." *IPSJ Trans. Computer Vision and Applications* 7 (2014): 10-17.
5. Li, Yang; Fu, Kun; Sun, Hao; Sun, Xian. 2018. "An Aircraft Detection Framework Based on Reinforcement Learning and Convolutional Neural Networks in Remote Sensing Images." *Remote Sens.* 10, no. 2: 243.
6. Xu T.B., Cheng G.L., Yang J. Fast Aircraft Detection Using End-to-End Fully Convolutional Network; Proceedings of the 2016 IEEE International Conference on Digital Signal Processing (DSP); Beijing, China. 16–18 October 2016
7. Zhang W., Sun X., Fu K., Wang C., Wang H. Object detection in high-resolution remote sensing images using rotation invariant parts based model. *IEEE Trans. Geosci. Remote Sens.* 2014;11:74–78. doi: 10.1109/LGRS.2013.2246538.
8. Ren S., He K., Girshick R.B. Object Detection Networks on Convolutional Feature Maps. *IEEE Trans. Pattern Anal. Mach. Intell.* 2017;39:1476–1481. doi: 10.1109/TPAMI.2016.2601099.
9. Dataset obtained from Kaggle - <https://www.kaggle.com/rhammell/planesnet/data>
10. information on mAP score - <https://towardsdatascience.com/breaking-down-mean-average-precision-map-ae462f623a52>